

Computer Vision – CNN Architectures

Fundamental CNN Architectures and best practices

- Given an application, several architecture design choices
 - #layers, #filters, kernel size, pooling, fully connected layers, regularizers etc.
- Look at competitions in related domains and start with these network designs, follow best practices
- Are there ways to reuse these open source networks trained on related domains without having to train from scratch?

ImageNet (ILSRC)

airplane



automobile




bird



cat



deer




dog



frog



horse



ship



truck



40 M images from 20k categories !

Top-5 error on 100k test images

Steel drum



Output:
Scale
T-shirt
Steel drum
Drumstick
Mud turtle

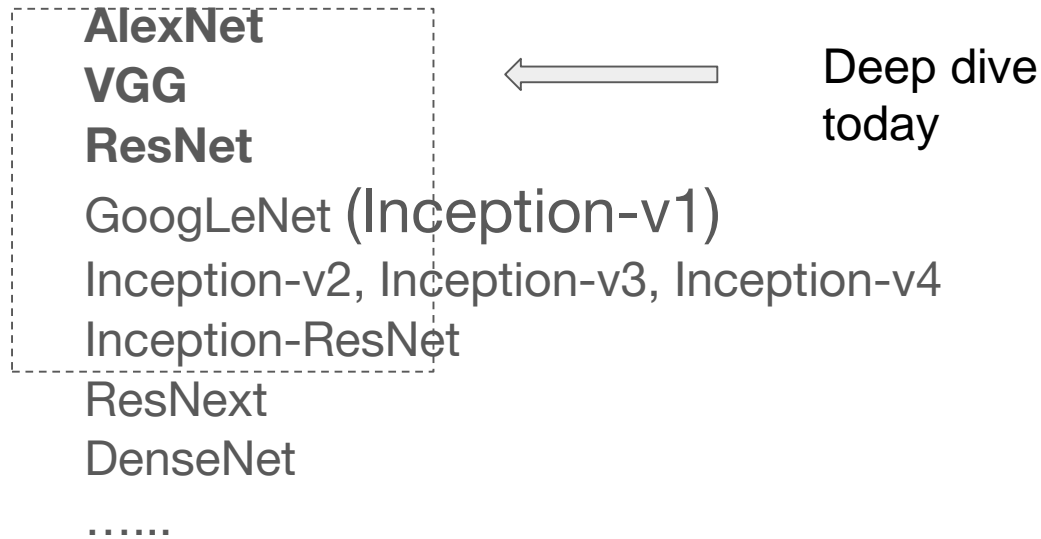


Output:
Scale
T-shirt
Giant panda
Drumstick
Mud turtle



$$\text{Error} = \frac{1}{100,000} \sum_{\substack{100,000 \\ \text{images}}} 1[\text{incorrect on image } i]$$

ImageNet trained Architectures and best practices



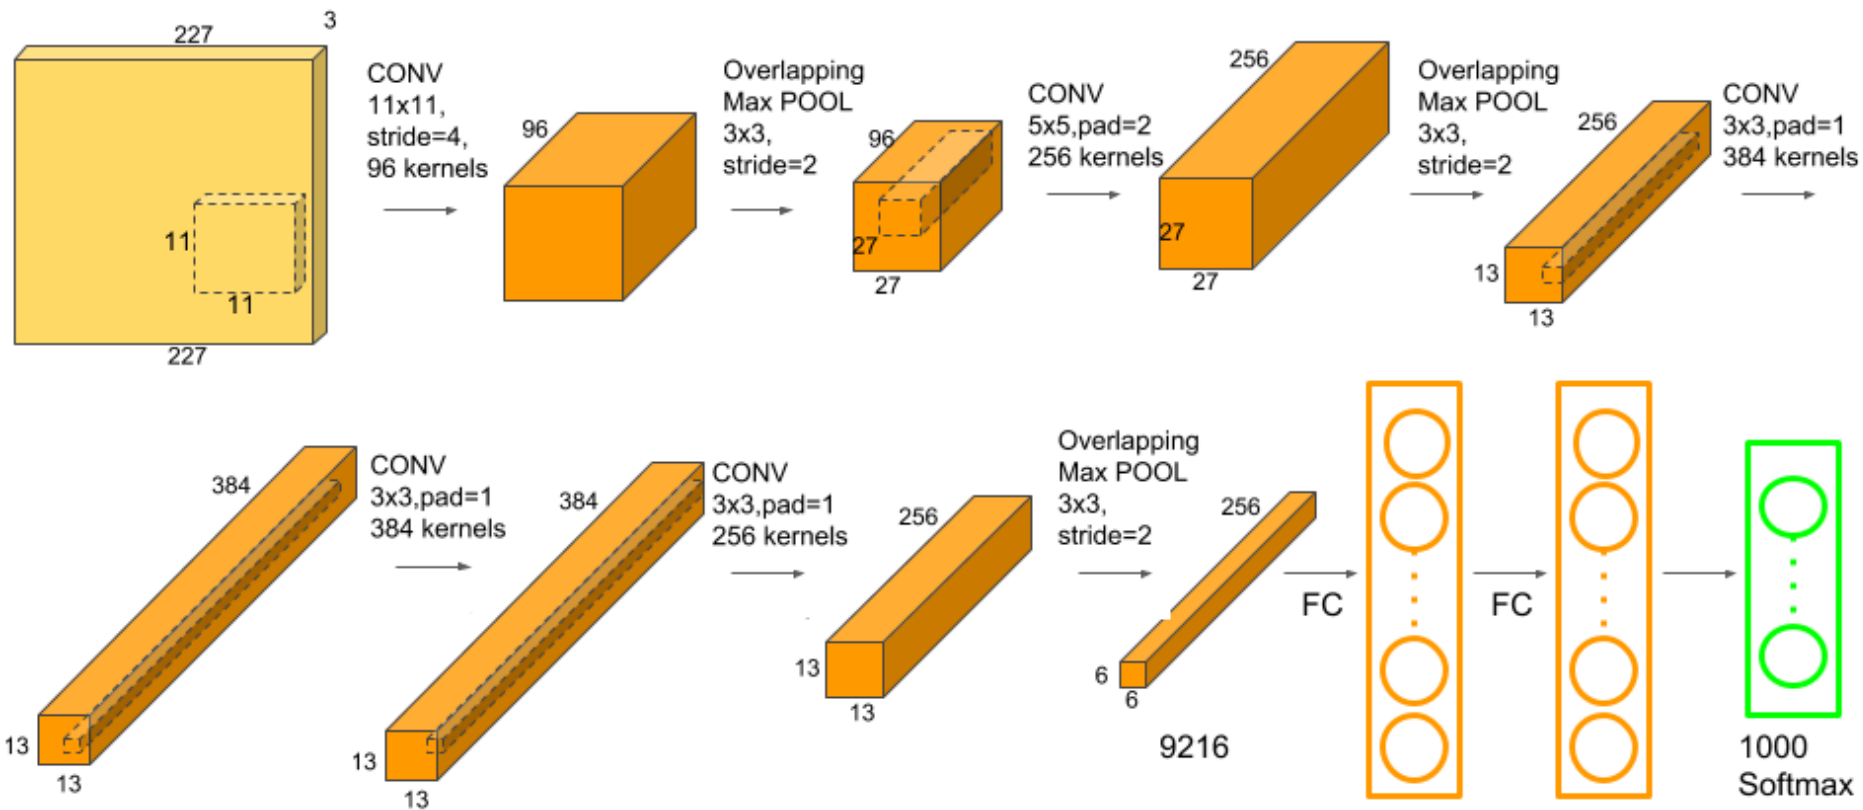
AlexNet(2012)

1. Around 2011, a good ILSVRC classification error rate was **25%**. In 2012, AlexNet achieved **16%**, a watershed moment!
2. Since then, the Computer Vision field has completely changed for one!
3. Compared to the state of the art DL architectures in 2012, AlexNet had a deep architecture (5 Conv layers, 3 Fully connected layers)

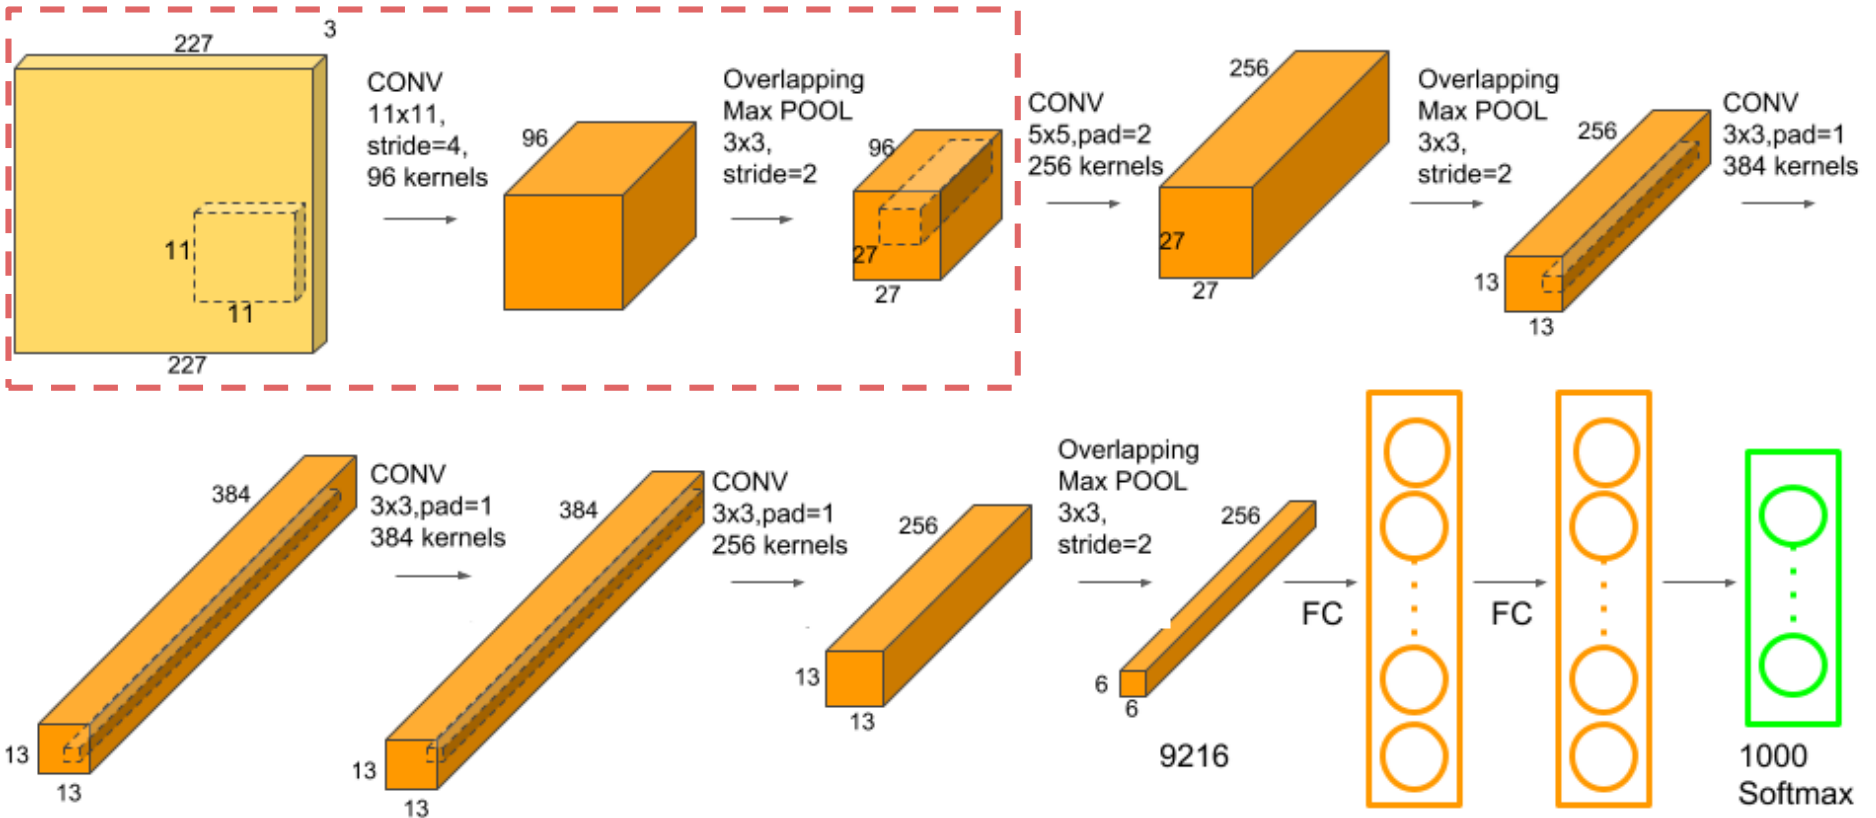
What made AlexNet successful?

1. AlexNet architecture
2. Deep dive block by block
3. Overlapping max pooling
4. ReLu
5. Dropouts
6. Cropping
7. Data Augmentation
8. Inference Augmentation

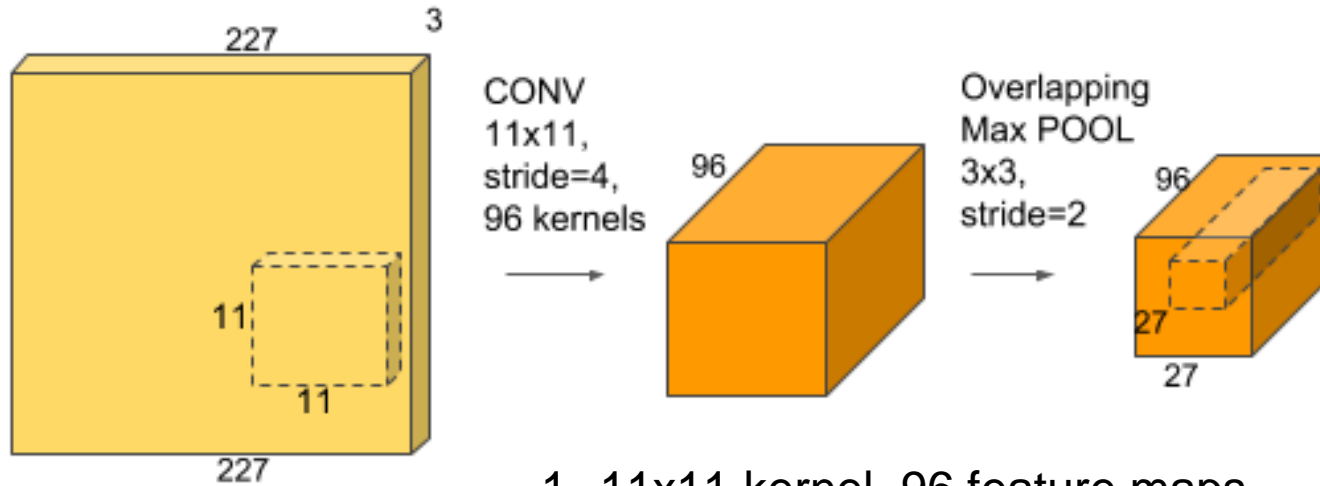
AlexNet (2012)



Lets Step in..

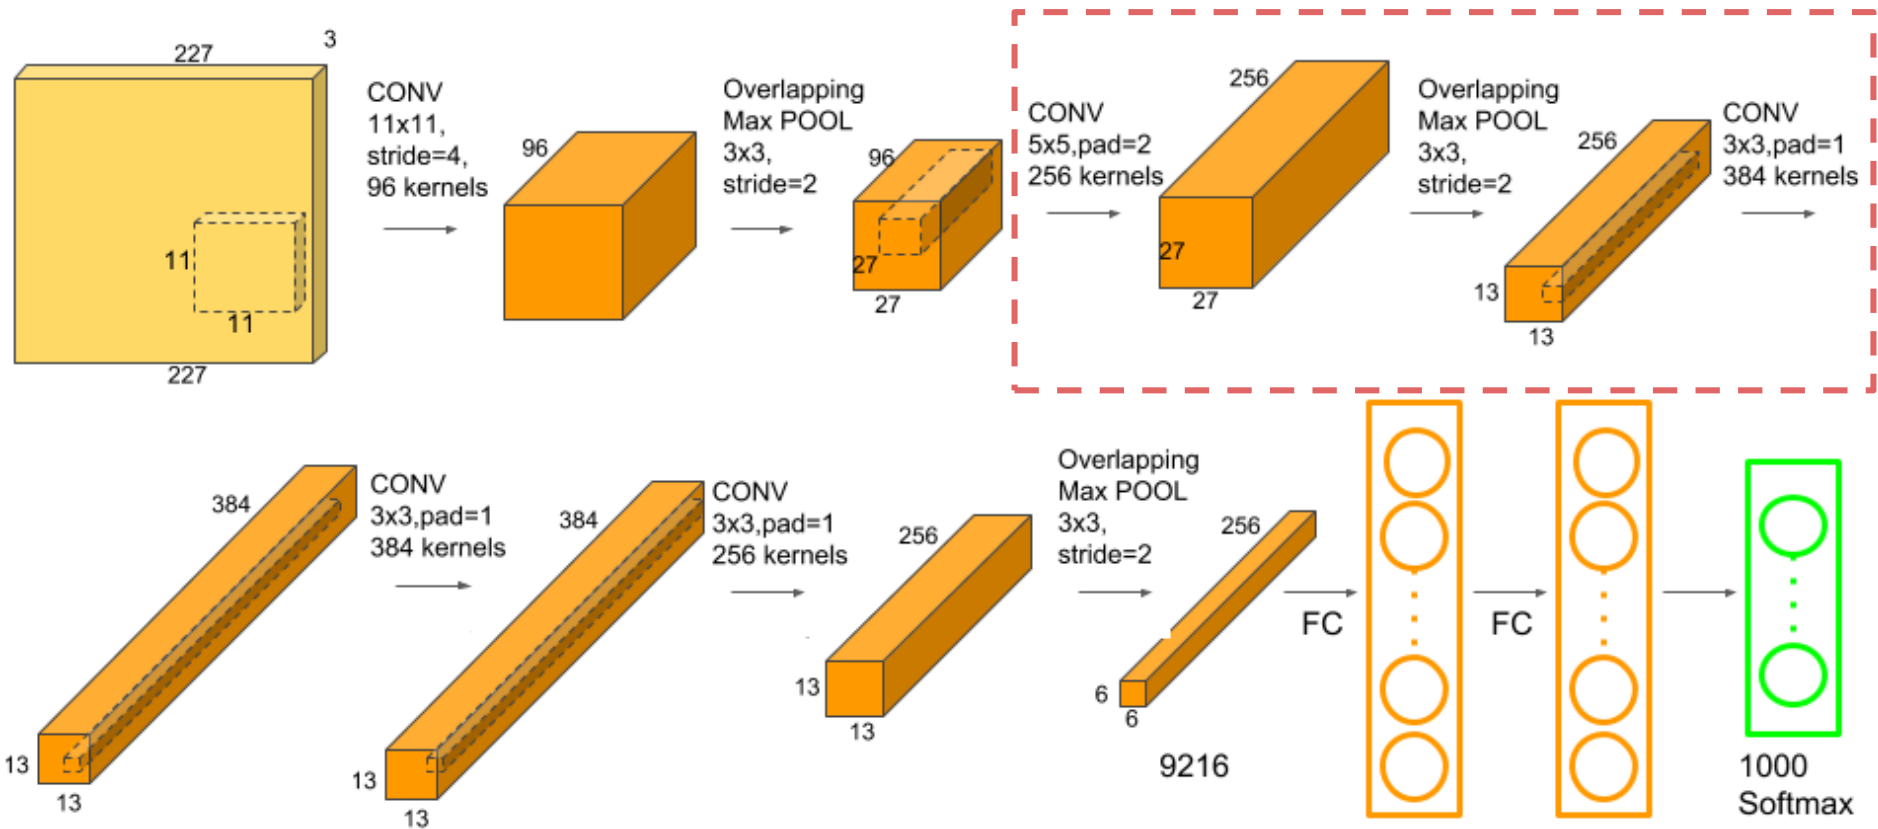


Lets go block by block - First Block





1. 11x11 kernel, 96 feature maps
2. Large Stride=4
3. Formula for output size - $(W+2P-F)/S+1$
4. Formula for no of parameters (ignore bias)- $M \times N \times k \times k$
5. Maxpool, 3x3, s=2

Lets Step in..



Second block






1. Flatten layer
2. FC layer
3. Softmax
4. #parameters in FC layers?

Overlapping Max Pooling

(3x3, stride 2)

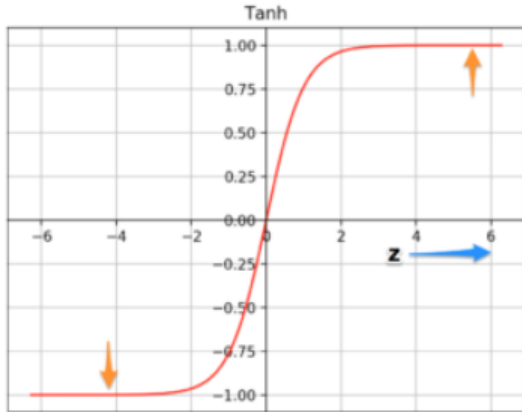
1	4	5	2	7
5	3	6	3	6
7	2	1	1	4
3	9	4	6	7
4	2	5	1	2




7	7
9	7

Moderate performance
gain reported by authors


ReLU instead of tanh



tanh




ReLU




Faster
convergence

Previously, VG was an issue networks couldn't go deeper, ReLU



(a) Standard Neural Net



(b) After applying dropout.


Regularization key due to huge #parameters in FC layer

Input Images

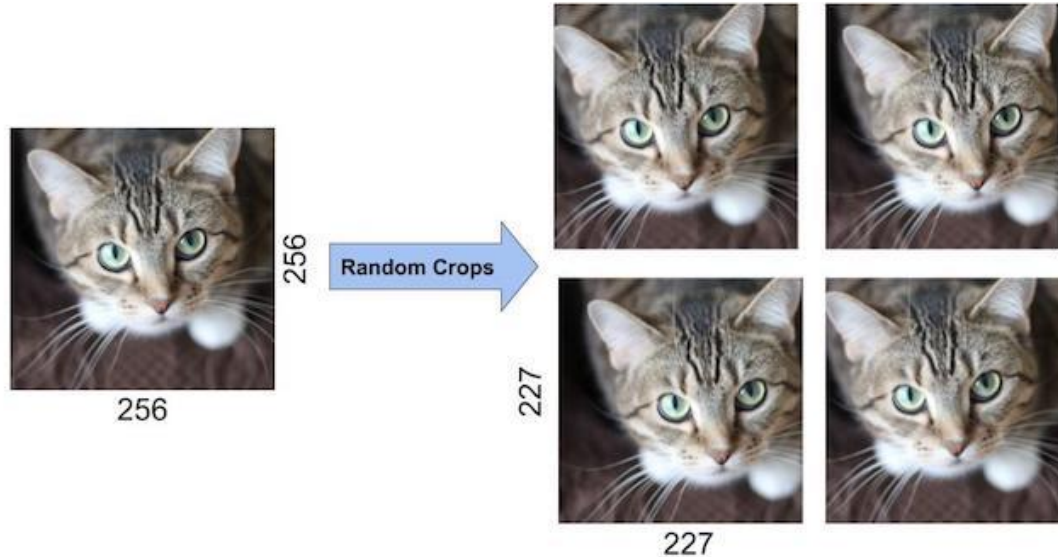


Resize smaller side to 256 and crop larger side to get 256x256
Get close to object of interest

Data augmentation




Horizontal flips



Random crop of 227x227 from 256x256 images

Inference Augmentation



Predict for each
image using
learnt model

Average
prediction

Is seen to improve accuracy moderately in many applications

**Test
image**

**Augmented
versions of test
image**

Summary

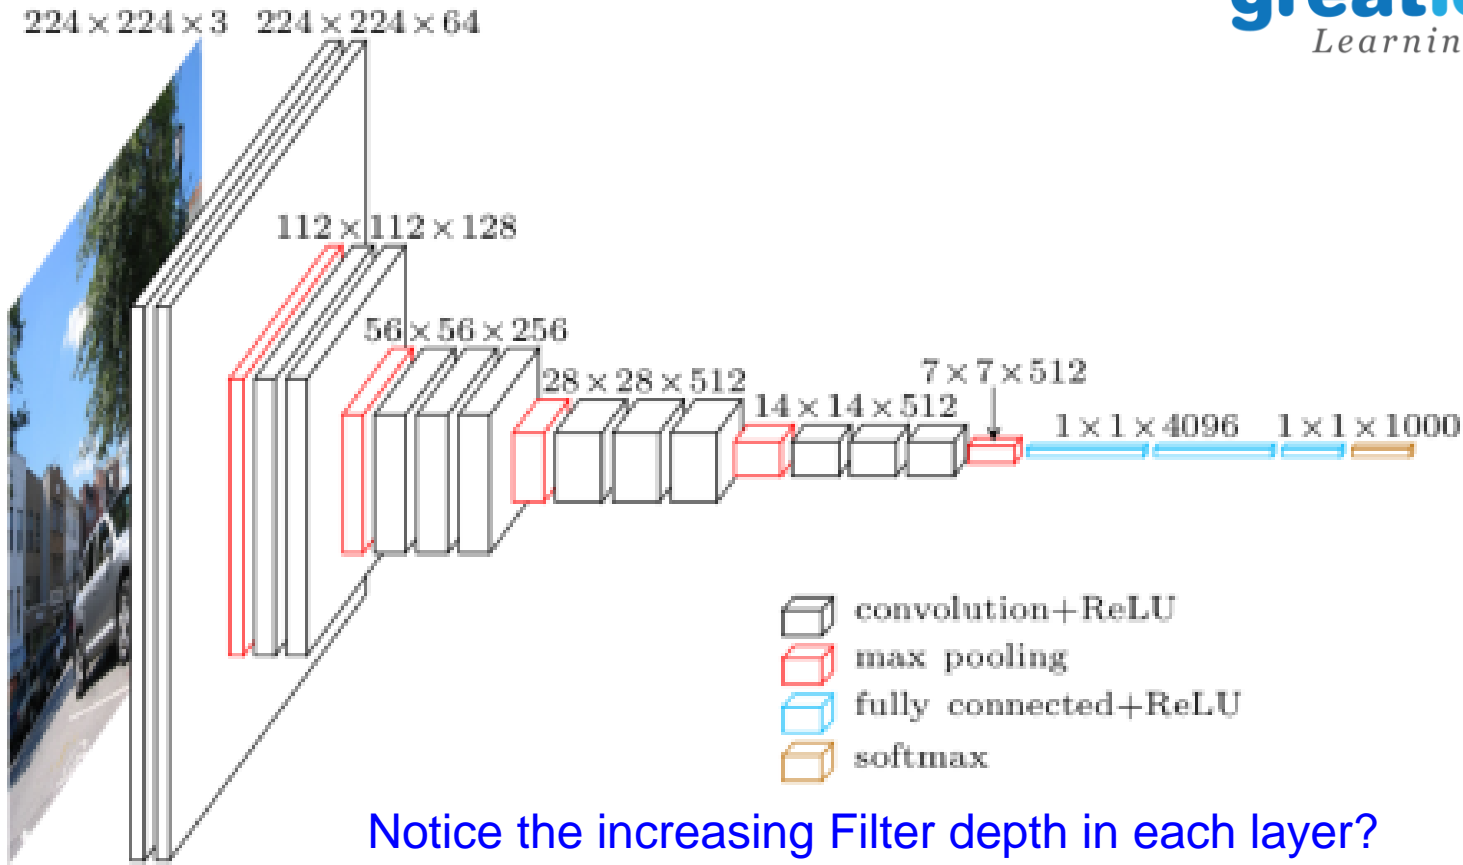
1. Trained the network on ImageNet data
2. Overlapping Maxpool
3. Used ReLU for the nonlinearity functions
4. Data augmentation: image translations, horizontal reflections, and patch extractions.
5. Inference/Test-time augmentation
6. Dropout in fully connected layers
7. Trained on two GTX 580 GPUs for five to six days

VGG (2014)

Another influential work is VGG which brought the ImageNet error down below **10%** (**7.3%** precisely)

Use of 3x3 filters is mimicked by most works today

Scale Augmentation at Train and Test time is another key addition



Key Points

Use of 3×3 Filters instead of large-size filters (such as 11×11 , 7×7)

Different VGG architectures

Increasing Filter depth

Multi-Scale Training/ Testing

Model Ensembling

Need for large filters and challenges

In images, non-local or wide range pixel interactions is important to capture

Thus a wide receptive field is important



Need
smaller
Receptive
field



Need
Larger
Receptive
field

Larger kernels (7x7, 11x11), Maxpooling are possibilities


With pooling, information loss is a risk

Larger kernels mean more parameters/compute


Remember: M- inputs of dimension $D \times D$, N - outputs, $K \times K$ kernels take $M \times N \times K \times K$ parameters and $M \times N \times K \times K \times D \times D$ operations

Use of 3x3 filters

Use of multiple layers of 3x3 filters instead of 1 layer of 5x5 or 7x7 or 11x11



two successive
3x3 convolutions



5x5 convolution

Receptive field

5x5 layer receptive field

Input



Feature
Map



5x5 filter and
nonlinear activation



Stacked 3x3 layer receptive field

Receptive field

Input



Feature Map 1



Feature Map 2



3x3 filter and
nonlinear activation

3x3 filter and
nonlinear activation

Same Receptive field and more non-linearity

Parameters/Computations

What is the number of parameters and receptive field in the following two cases

Input channels = 32

conv-32, k=3x3, s=1,'relu'

conv-32, k=3x3, s=1,'relu'

Input channels = 32

conv-32, k=5x5, s=1,'relu'

Parameters/Computations

What is the number of parameters and receptive field in the following two cases

Input channels = 32

conv-32, k=3x3, s=1,'relu'

conv-32, k=3x3, s=1,'relu'

$$32 \times 32 \times 3 \times 3 + 32 \times 32 \times 3 \times 3 = 32 \times 32 \times 18$$

Input channels = 32

conv-32, k=5x5, s=1,'relu'

$$32 \times 32 \times 5 \times 5 = 32 \times 32 \times 25$$

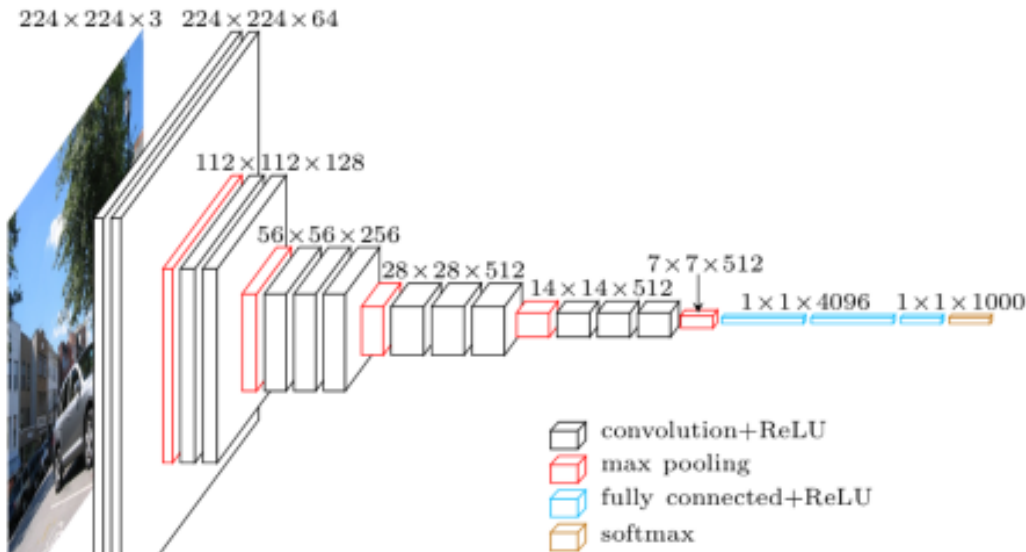
Remember both have same receptive field of 5 x 5 !

Different VGG architectures

Architectures
used in the
VGG work

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					

Increasing Filters with Depth




Initial layer
depth

ution, less

Upper layers encode high-level info, less spatial resolution, more depth

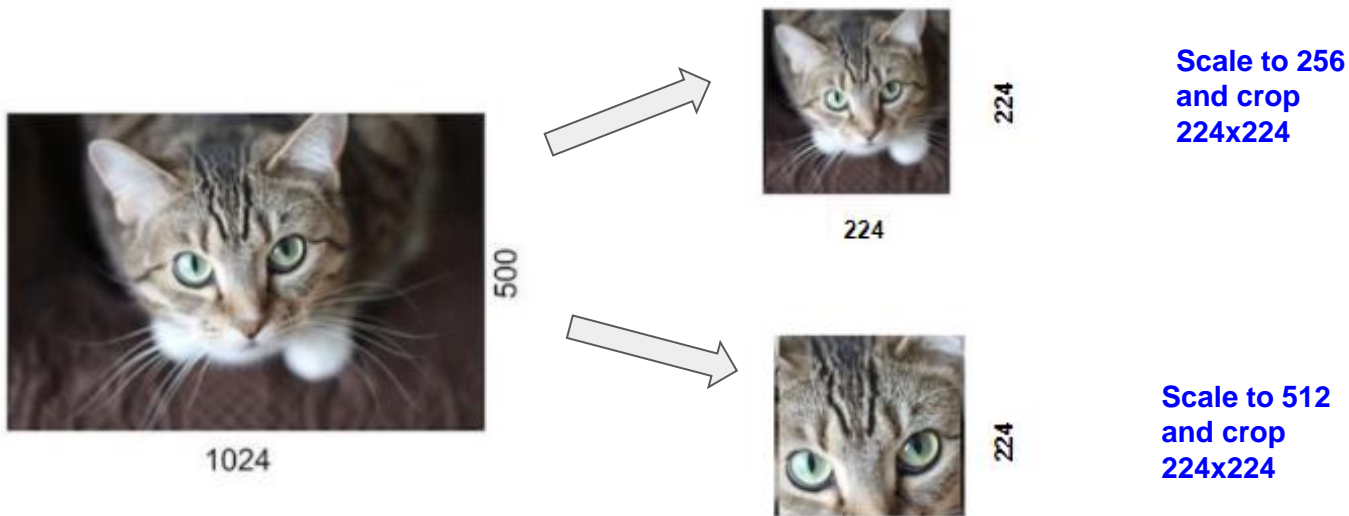
Maintain information content with decreasing spatial resolution

Rectangular Input Images: Cropping



Resize smaller side to 256 and crop at center to get 224x224

Multi-scale augmentation at train/test time



Resize smaller side to multiple scales in $[256, 512]$ and crop to get 224x224


Multi-scale augmentation at train/test time



Same image gives
different scaled and
cropped/shifted
versions

Model ensembling

Average prediction probabilities from multiple models (VGG-16, VGG-19)



Handle Overfitting issues,
higher accuracy

Summary

1. The use of only 3x3 sized filters as against AlexNet's 11x11 filters in the first layer.
2. Increasing filters with depth
3. Used scale variation as one data augmentation technique during training and testing.
4. Model ensembling for best results
5. The top-5 test error on ImageNet was 7.3%

Residual Networks


Were the first to train really deep networks (150 layers, 1000 layers)

Imagenet error rate down to **3.57%** from **7.32 %** (VGG)

Very key idea of Residual connections

Deep Residual Networks

- Neural Networks with just 1 hidden layer are universal approximators
- Efficient representation is important for managing computational requirements, robust learning and preventing overfitting
- An important element of representation is depth of the network
- The benefit of depth has been successfully demonstrated previously in AlexNet, VGG



Advantages of greater Depth

- Representation complexity grows exponentially w.r.t hidden units compared to shallow networks
- Thus for same number of parameters, Deep networks allow for more complex representation
- Deep CNN networks with small filters (e.g. 3x3, 1x1) have lesser parameters/faster compute for same receptive field

Challenges of Deep Networks-Hard to train

- Vanishing gradients issue
 - RELU alleviates this issue to some extent
- Degradation in training
 - Increased non-convexity, harder to train
 - Simple maps like the identity map hard to converge




Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

Some Math Examples

$$P(x) = f(g(x))$$

$$P'(x) = f'(g(x))g'(x)$$

$$P''(x) = f''(g(x))(g'(x))^2 + f'(g(x))g''(x)$$

See the Vanishing Gradient problem?

Is the composition of 2 convex functions convex?

Some Math Examples

$$P(x) = f(g(x))$$

$$P'(x) = f'(g(x))g'(x)$$

$$P''(x) = f''(g(x))(g'(x))^2 + f'(g(x))g''(x)$$

Is the composition of 2 convex functions convex?

Not in general. e.g. e^{-x^2}

What if we use residuals ?

Convex i.e. the 2nd derivative is positive definitely if f is monotonic.


$$P(x) = f(x + g(x)) + x + g(x)$$

We define f and g
as residuals

$$P(x) = (f'(x + g(x)) + 1)(g'(x) + 1)$$

The 1st derivative is
better conditioned


Residual Block: Skip connection



In I

Figure 2. Residual learning: a building block.

Residual Block: Skip connection



As

Figure 2. Residual learning: a building block.

well behaved

Residual Block: Skip connection


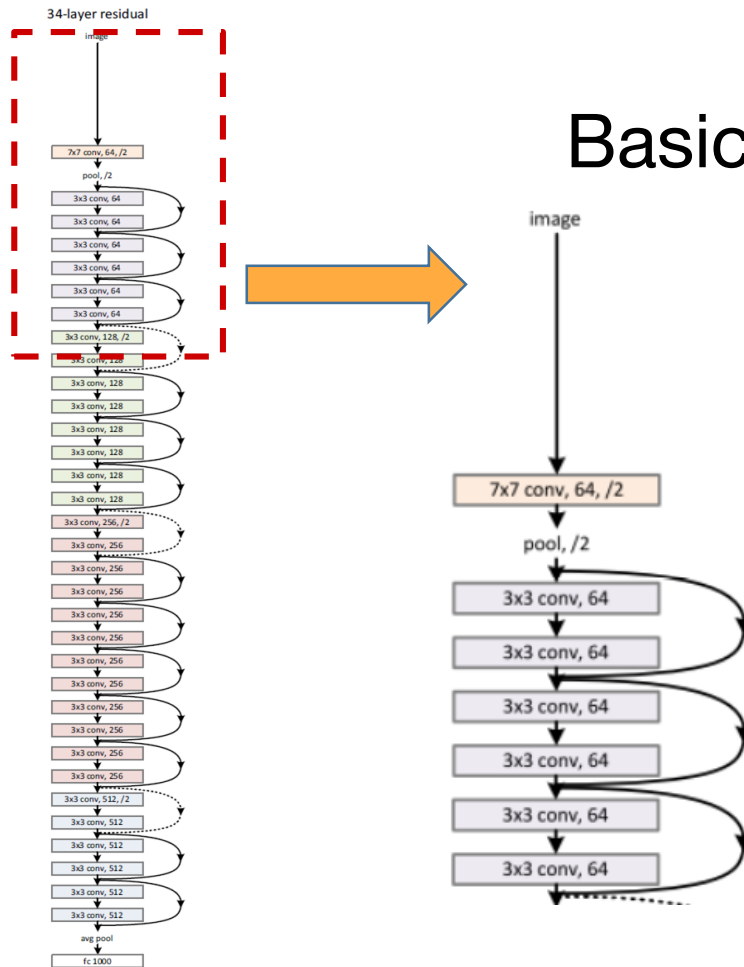


Figure 2. Residual learning: a building block.



Typically the Residuals are small...

Thus, by stacking more layers, worst case is, we learn the identity. Earlier, the entire layer would collapse!

Basic Architecture- Resnet 34





34-layer residual





greatlearning
Learning for Life

Basic Architecture- Resnet 34

Basic Architecture- Resnet 34



Improved Training and Test Accuracy



Summary


- Skip connections for training very deep networks
- Scale/horizontal flip data augmentation
- Batch normalization
- Dropout not used
- Fully convolutional output
- Multi-crop/multi-scale prediction and averaging testing
- Imagenet error rate down to **3.57%** from **7.32 %** (VGG)

GoogLeNet/Inception-v1

Topics we will look at:

1. The 1×1 Convolution
2. Inception Module
3. Global Average Pooling
4. Overall Architecture
5. Auxiliary Classifiers for Training
6. Testing Details


1x1 Convolution filters



Each
by n


1x1 filters

ture maps followed



Easy way to get Feature reduction/increase,
additional non-linearity


Reduction of parameters



#Parameters - $48 \times 480 \times 5 \times 5 = \mathbf{0.5\ M}$


#OPs - $14 \times 14 \times 480 \times 5 \times 5 \times 48 = \mathbf{113M}$

Without the Use of 1×1 Convolution



With the Use of 1×1 Convolution


#OPs for the below network - ?



#Parameters - $48 \times 480 \times 5 \times 5 = \mathbf{0.5 \text{ M}}$

#OPs - $14 \times 14 \times 480 \times 5 \times 5 \times 48 = \mathbf{113 \text{ M}}$

Without the Use of 1×1 Convolution




With the Use of 1×1 Convolution

#OPs - $14 \times 14 \times 480 \times 16 + 14 \times 14 \times 16 \times 5 \times 5 \times 48 = \mathbf{5.3 \text{ M}}$

Possible ways to derive the Output feature map

The Object is identifiable by just a linear combination of input features/channels



Objects in an image is small requiring small kernel size



224

Objects could be big

224

larger sized kernel



224

224

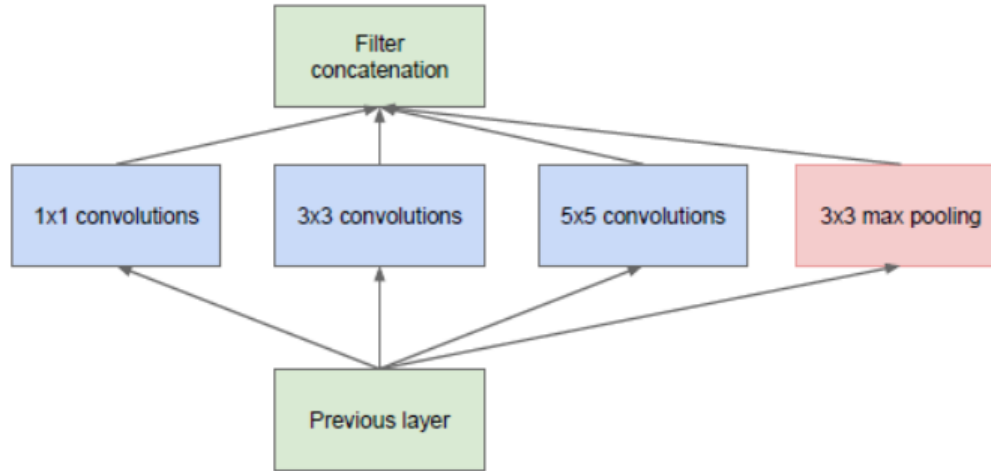
Do we need to focus on a lower resolution or same resolution for classification

Do we need Pooling or not?

Thus, at every layer, there is a design choice on a) linear combination of input maps b) size of kernels c) Whether or not to do Pooling

Can we use data/optimization to choose on what is important for a layer ?

Inception Block !



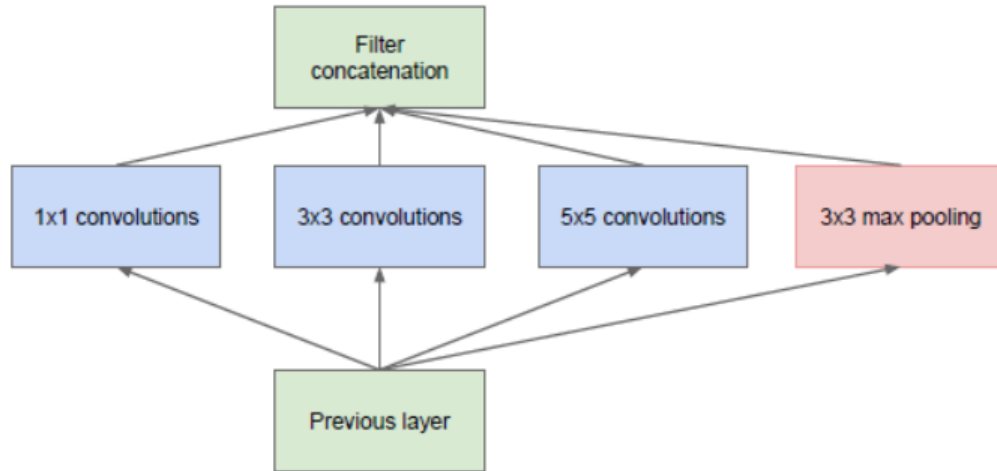
(a) Inception module, naïve version

Offset

3x3 or 5x5 or combinations c) Whether or not to do Pooling

size of kernels

Too Many parameters and Expensive

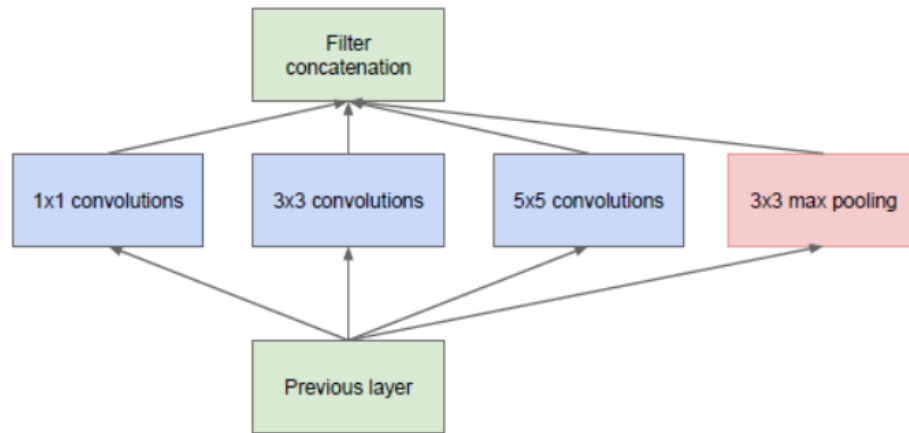


(a) Inception module, naïve version

Say input dim = $256 \times 256 \times 4$. Can you compute the #parameters and #OPs for Output = 4 features for each path

e.g. 1x1 conv - #Param: 4×4 #OPs: $256 \times 256 \times 4 \times 4$

Too Many parameters and Expensive



(a) Inception module, naïve version

Say input
Output =

1x1 conv - #Param: 32×32 #OPs: $256 \times 256 \times 32 \times 32$

So What to do?

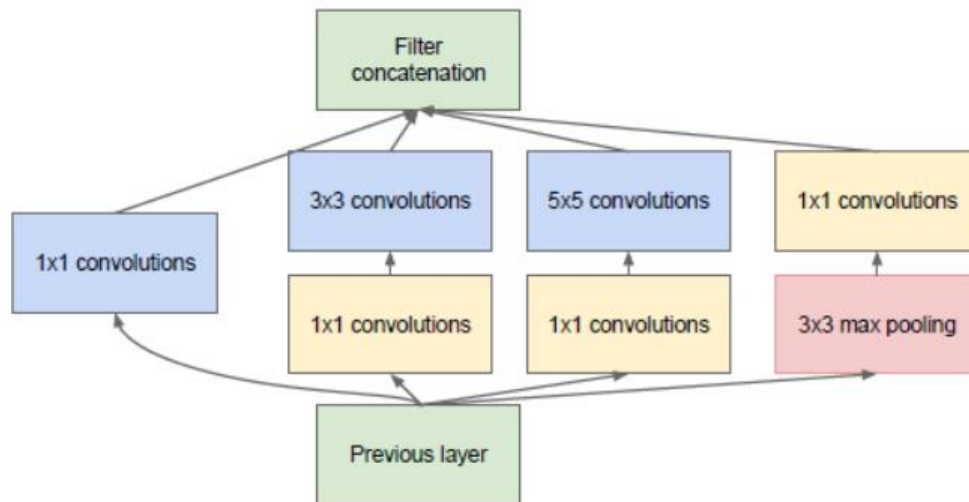
3x3 conv - #Param: $32 \times 32 \times 3 \times 3$ #OPs: $256 \times 256 \times 32 \times 32 \times 3 \times 3$

5x5 conv - #Param: $32 \times 32 \times 5 \times 5$ #OPs: $256 \times 256 \times 32 \times 32 \times 5 \times 5$

3x3 pool - #Param: 0 #OPs: $128 \times 128 \times 32 \times 3 \times 3$


ers and #OPs for

Use 1x1 conv to reduce parameters and speed



(b) Inception module with dimensionality reduction

Global Average Pooling





Fully Connected Layer VS Global Average Pooling

FC parameters: $7 \times 7 \times 1024 \times 1024 = 501,177,600$



GAP parameters: **0**

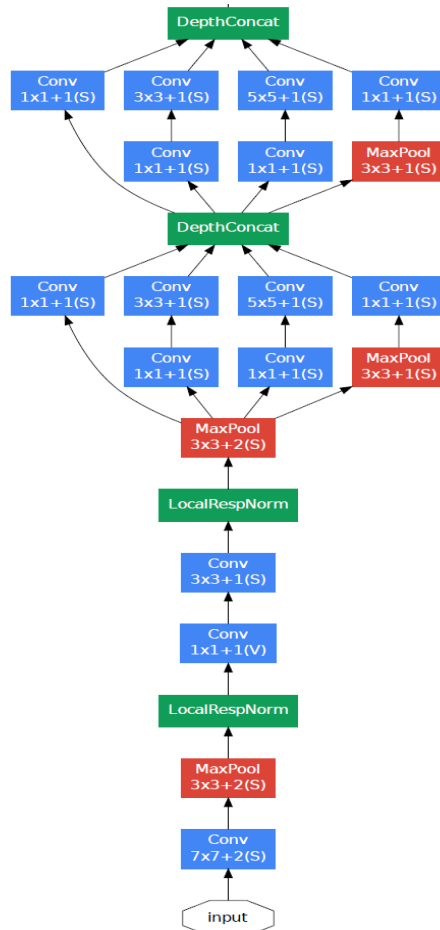
Much less parameters using GAP, less overfitting!

Overall Architecture





Overall Architecture






Overall Architecture



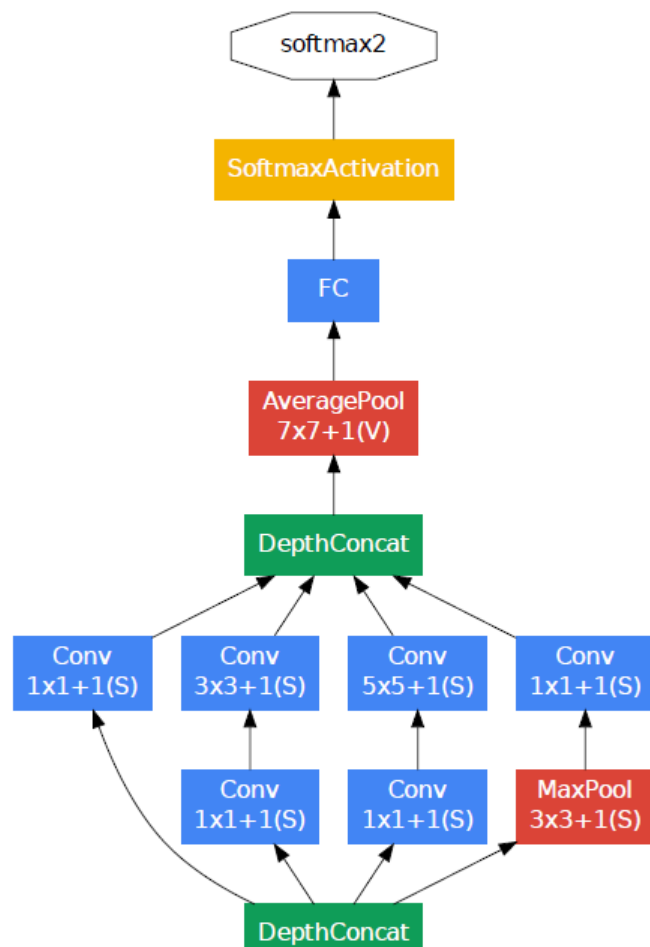


Use this to
avoid VG effect
in middle layers

Overall Architecture









Summary

- Used 9 Inception modules, 100 layers in total
- No use of fully connected layers. This saves a huge number of parameters.
- Uses 12x fewer parameters than AlexNet.
- During testing, multiple crops of the same image were created, fed into the network, and the softmax probabilities were averaged to give us the final solution.
- Trained on “a few high-end GPUs **within a week**”.
- Imagenet Top-5 error rate down to **6.66%** from **7.32 %** (VGG)

State of Art CNN architectures

Performance trends (ImageNet (<https://en.wikipedia.org/wiki/ImageNet>))



An Analysis of Deep Neural Network Models for Practical Applications, 2017.